

Kubernetes Fundamentals

(Greek (Κυβερνήτης) -> kee-ver-NEE-tees)



Kubernetes - K8s



Numeronym, “a number-based word”

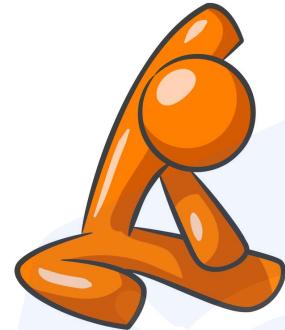
Examples:

- Phonetic:
 - K9 = Canine, “Police Dogs”
- Number of letters omitted:
 - l10n = localization
 - i18n = internationalization
 - K8s = Kubernetes

Agenda



- Introduction to Kubernetes
- Namespaces, Pods, and Deployments
- Service Discovery and Load-balancing
- Rolling out updates
- ConfigMaps and Secrets
- Persistent Storage
- Networking and Ingress
- Helm Package Manager
- Kubernetes Internals



Agenda



Why Kubernetes?

Agenda



Docker is awesome
but
Docker isn't enough

Agenda



Docker is awesome
for packaging

Kubernetes is awesome
for orchestration

(Some) Orchestration with Docker Compose!



docker-compose.yaml

```
version: '3'
services:
  nginx:
    image: "nginx"
    ports:
      - "80:80"
  mysql:
    image: "mysql"
    ports:
      - "3306:3306"
```

eficode

Kubernetes Fundamentals www.eficode.com

More Problems



- Multiple application stacks
- Large deployments
- Deployment / Scheduling / Updates / Rollback
- Replication / Flexible scalability
- Resilience / Availability
- Storage
- Ingress / Load balancing
- Security

eficode

Kubernetes Fundamentals www.eficode.com

We need an orchestration tool!



- Availability - scale as defined in your desired state
- Resilience - if a container exits/dies, a new one can be automatically created
- Storage - Local, NFS, iSCSI, “cloud”, etc.
- Deployments - with patterns like Canary, A/B testing, Blue/Green, etc.
- Scheduling - with Resource Limitations
- Updates - with Rolling Updates
- Networking and Cluster DNS
- Service Discovery
- Ingress routing and load balancing across replicas
- Security - with RBAC and network policies

Why Kubernetes?



- Industry Standard for Container Orchestration
 - **Broad industry support** for managed solutions
 - Backed by Cloud Native Computing Foundation (CNCF)
- Strong Innovation in the K8s Area
 - **Commoditizing the container-based application infrastructure.** See [CNCF landscape](#)
 - Applications are built to run in K8s (+600 Official Helm charts)
- Based on experience from existing solutions (Google Borg, Docker)
 - Opinionated on how you **model** your container application
 - Not opinionated on the implementation of the K8s platform
 - Almost every component has a number of alternatives..

Agenda



How Kubernetes?

eficode

Kubernetes Fundamentals www.eficode.com

The Short Description of Kubernetes



Application blueprint

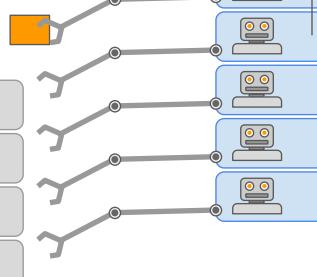


*"Build and maintain
my application
according to this
blueprint"*

Blueprint storage



Container nodes



Controllers

Scheduling, resource mgr

Container instance manager

Storage manager

Security compliance mgr

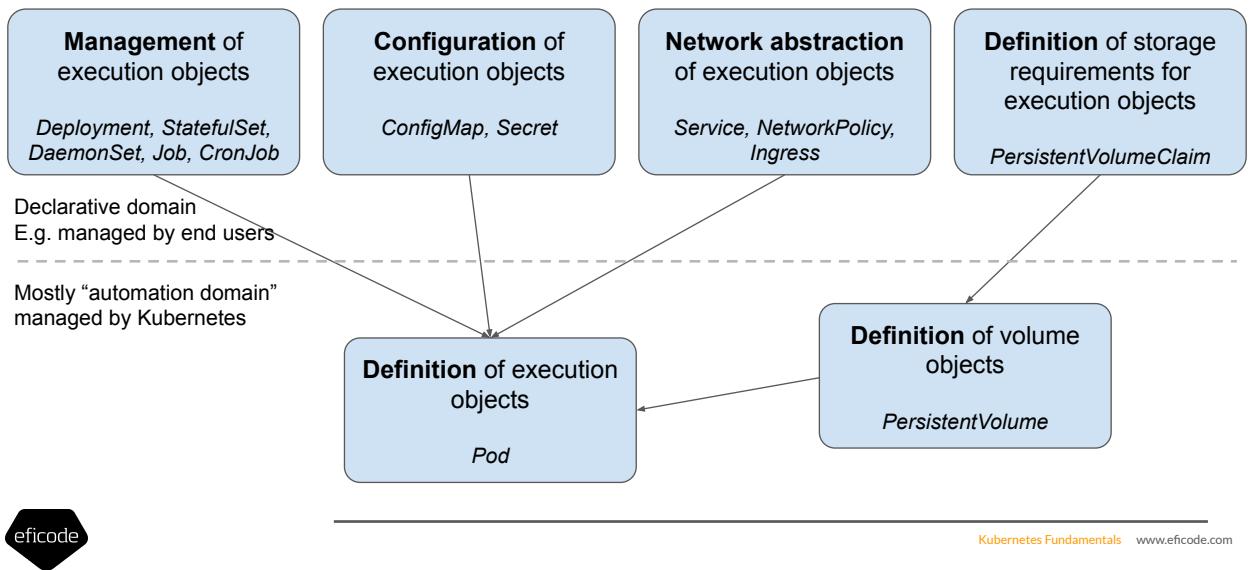
Network manager

... etc

eficode

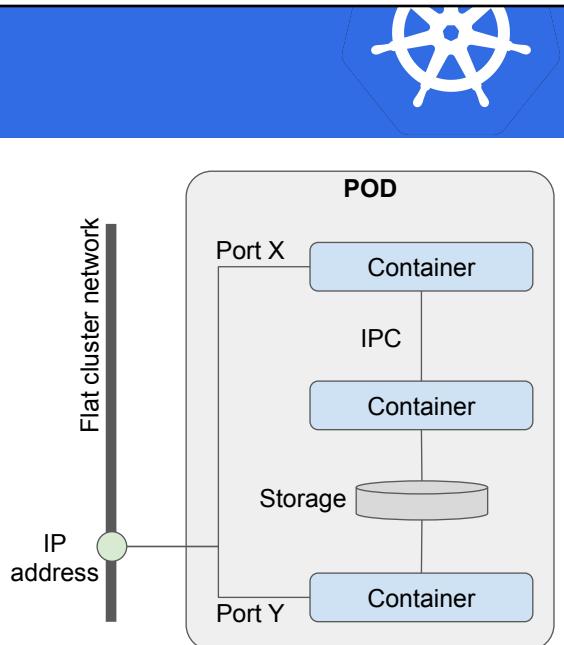
Kubernetes Fundamentals www.eficode.com

What are the Kubernetes Lego Bricks?



The POD Model

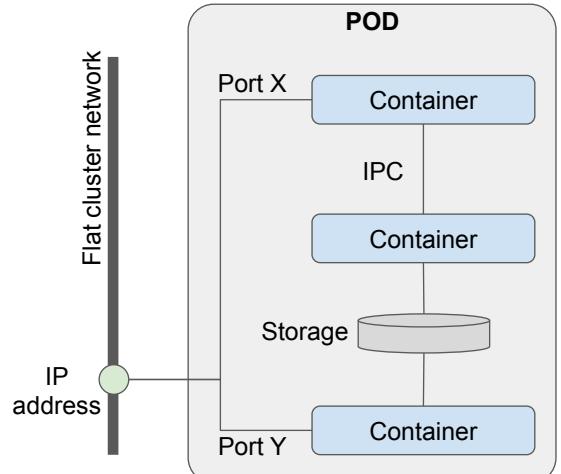
- Smallest object in K8s is a POD
- Consists of one or more containers
- IP addressable with ports (like a VM)



The POD Model



- Smallest object in K8s is a POD
 - Replicated as a unit.
 - Has a life-cycle.
 - Is scheduled to worker nodes.



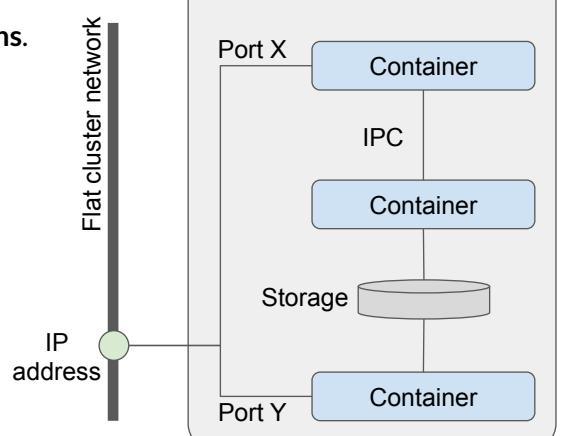
eficode

Kubernetes Fundamentals www.eficode.com

The POD Model



- Consists of one or more containers
 - Grouping enables strong **use-case patterns**.
(To be continued.)
 - NB: Only put highly related containers together in a POD**
 - Scaled, live, die together.**



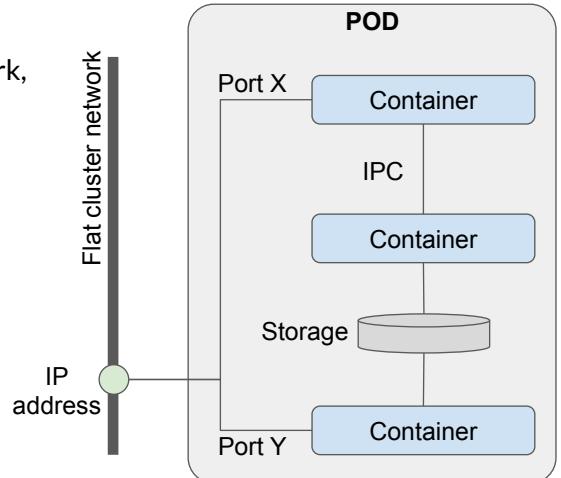
eficode

Kubernetes Fundamentals www.eficode.com

The POD Model



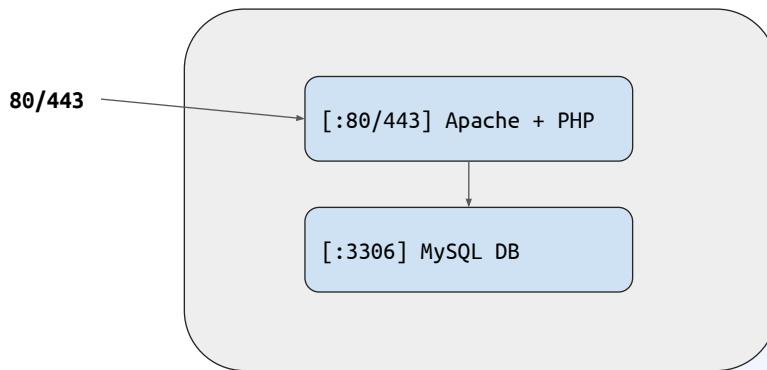
- IP addressable with ports (like a VM)
 - PODs always see network as a flat network, no NAT.



Kubernetes Fundamentals www.eficode.com

eficode

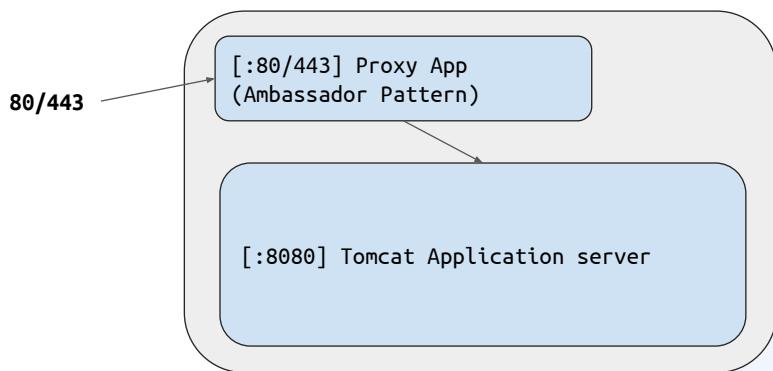
(Bad) Example of multi-container pod



Kubernetes Fundamentals www.eficode.com

eficode

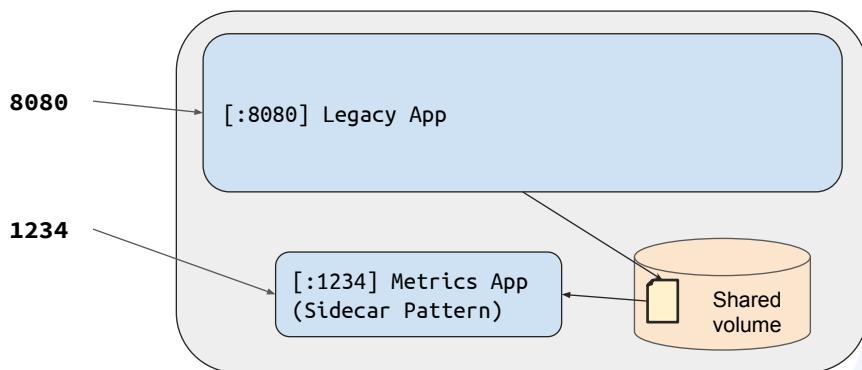
(Good) Example of multi-container pod



eficode

Kubernetes Fundamentals www.eficode.com

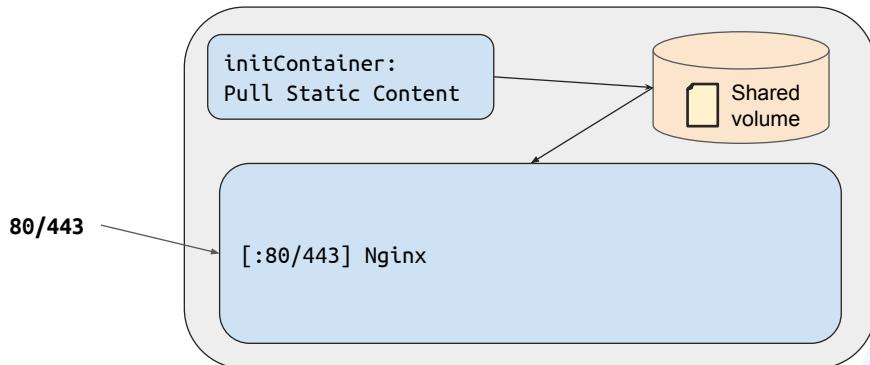
(Good) Example of multi-container pod



eficode

Kubernetes Fundamentals www.eficode.com

(Good) Example of multi-container pod



eficode

Kubernetes Fundamentals www.eficode.com

Namespaces



"Virtual clusters" in Kubernetes

Cluster

Namespace A
Test environment

Pod A



Namespace B
Production

Pod B



Namespace C
Developer environment

Pod C



eficode

Kubernetes Fundamentals www.eficode.com

Workshop Setup



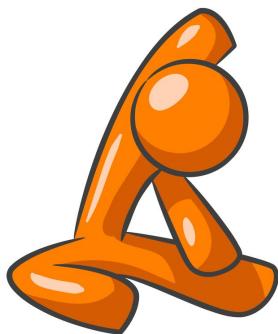
- A shared Kubernetes cluster on Google Cloud Platform
- Run exercises from your VM
 - \$ kubectl get nodes -o wide
- Each trainee have a **private** namespace
- A very good tool: **praqma/network-multitool**

Exercise 1 - first half



<https://github.com/eficode-academy/kubernetes-katas>

01-pods-deployments.md



NB: Do the optional 1.1.1, but **STOP** the exercise before section 1.2!

- Create a pod in your namespace
- Try accessing the pod with your web browser

Interacting with Kubernetes



- Command line: kubectl
 - Imperative (create, delete)
 - Declarative (apply -f file.yaml)
- Point and click: K8s Dashboard
 - Very imperative
- Templating
 - Very Declarative
 - Helm, Kustomize
- Helm orchestration
 - Helmfile, Helmsman

The screenshot shows the Kubernetes Dashboard interface. The left sidebar includes links for Overview, Nodes, Persistent Volumes, Roles, Storage Classes, Namespaces (All namespaces), Workloads, Daemon Sets, Deployments, Jobs, Pods, Replication Sets, Stateful Sets, Discovery and Load Balancing, Ingresses, Services, Endpoints, Config Maps, Persistent Volume Claims, and Secrets. The main content area displays four sections: Daemon Sets, Deployments, and Pods, each listing their respective objects with columns for Name, Namespace, Labels, Pods, Age, and Images.

Kubernetes Fundamentals www.eficode.com

eficode

Interacting with Kubernetes (2)



Declarative - the “blueprint” approach vs. imperative - the “hacking” approach

```
apiVersion: v1
kind: Pod
metadata:
  name: frontend
  labels:
    app: frontend
spec:
  containers:
    - name: webserver
      image: nginx
      ports:
        - containerPort: 80
    - name: myapp
      image: nodeapp
```

Run an nginx Deployment and expose it to the world:

```
docker run -d --restart=always -e DOMAIN=cluster \
--name nginx-app -p 80:80 nginx
```

On Kubernetes:

```
kubectl run --image=nginx nginx-app --port=80 \
--env="DOMAIN=cluster"
```

Kubernetes Fundamentals www.eficode.com

eficode

Defining Resources - POD YAML Example



Kubernetes resources are typically defined using YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: my-webserver
spec:
  containers:
    - name: nginx-container
      image: nginx
      ports:
        - containerPort: 80
```

A Pod spec.

Looks daunting.

Is not.

Defining Resources - POD YAML Example



Kubernetes resources are typically defined using YAML

```
apiVersion:
kind:
metadata:
  name:

spec:
  ...
```

All resources have

- apiVersion
- kind
- metadata
 - name
- spec

Defining Resources - POD YAML Example



Kubernetes resources are typically defined using YAML

```
apiVersion: v1
kind:
metadata:
  name:

spec:
  ...
```

apiVersion

“What API is this resource part of?”

- Beta? (might change)
- v1? (very unlikely to change)
- networking.istio.io/v1? (3rd party)

Defining Resources - POD YAML Example



Kubernetes resources are typically defined using YAML

```
apiVersion: v1
kind: Pod
metadata:
  name:

spec:
  ...
```

kind

“What is this resource?”

- Pod
- Deployment
- Service
- etc.

Defining Resources - POD YAML Example



Kubernetes resources are typically defined using YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: my-webserver
spec:
  ...
```

metadata

nb: always has a name!

“Extra information about this resource?”

- labels (n-dimensional selectors)
- annotations (arbitrary KV data)

Defining Resources - POD YAML Example



Kubernetes resources are typically defined using YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: my-webserver
spec:
  containers:
    - name: nginx-container
      image: nginx
      ports:
        - containerPort: 80
```

spec

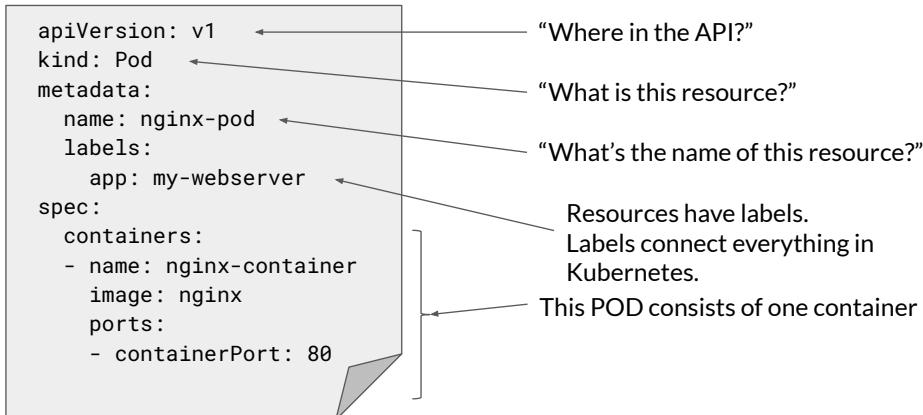
“How’s this resource configured?” ... very specific

- Pods have containers, volumesMounts
- Deployments have a template.spec
(= Pod.spec), number of replicas, etc.
- Services have network settings...

Defining Resources - POD YAML Example



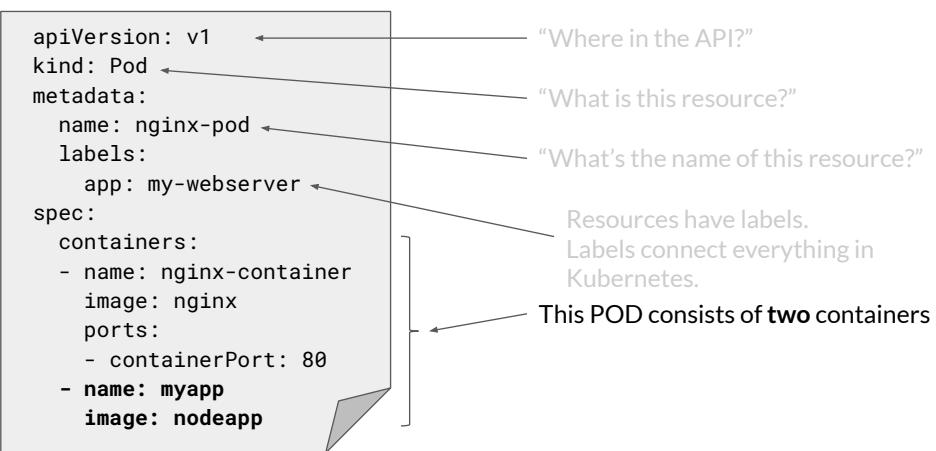
Kubernetes resources are typically defined using YAML



Defining Resources - POD YAML Example



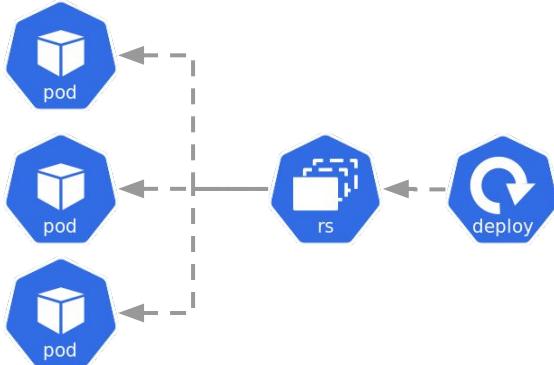
Kubernetes resources are typically defined using YAML



Managing Deployments in Kubernetes

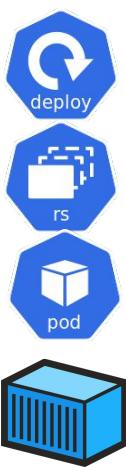


How do we scale our pods?



- Replica set
- Deployment

Managing Deployments in Kubernetes



Deployments manage the release of new versions of applications

ReplicaSets define how many instances should be running

Pods are the k8s scheduling primitive

Containers provide a runtime for a process

Use cases:

- Create
- Update
- Rollback
- Scale
- Pause/Resume

Labels Connects Things in Kubernetes



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: my-webserver
    environment:testing
    partition:customer-a
spec:
  containers:
    - name: nginx-container
      image: nginx
      ports:
        - containerPort: 80
```

- Everything can have labels.
- Labels are arbitrary key-value pairs
- The glue between most resources

Using multiple labels allows for an N-dimensional grouping of resources

eficode

Kubernetes Fundamentals www.eficode.com

Example 2 - A Deployment



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: my-webserver
spec:
  selector:
    matchLabels:
      app: my-webserver
  replicas: 4
  template:
    metadata:
      labels:
        app: my-webserver
    spec:
      containers:
        - name: nginx-container
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

“apps/v1” !? - look in the “Reference”

Higher level abstractions (controller) allowing Kubernetes to do more automation

We want four replicas - Kubernetes will deploy and manage these for us

Based on this POD template

eficode

Kubernetes Fundamentals www.eficode.com

Example 2 - A Deployment



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: my-webserver
spec:
  selector:
    matchLabels:
      app: my-webserver
  replicas: 4
  template:
    metadata:
      labels:
        app: my-webserver
    spec:
      containers:
        - name: nginx-container
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

"Deployments have a template.spec (= Pod.spec).."

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: my-webserver
spec:
  containers:
    - name: nginx-container
      image: nginx:1.7.9
      ports:
        - containerPort: 80
```

Example 2 - A Deployment

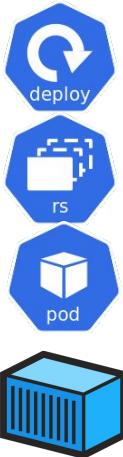


```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: my-webserver
spec:
  selector:
    matchLabels:
      app: my-webserver
  replicas: 4
  template:
    metadata:
      labels:
        app: my-webserver
    spec:
      containers:
        - name: nginx-container
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Expose the app as a service
on the cluster network using
labels for selection

```
apiVersion: v1
kind: Service
Metadata:
  name: nginx-service
  labels:
    app: my-webserver
spec:
  ports:
    - port: 80
      targetPort: 80
  selector:
    app: my-webserver
```

Managing Deployments in Kubernetes



Deployments manage the release of new versions of applications

ReplicaSets define how many instances should be running

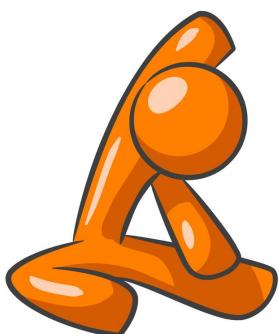
Pods are the k8s scheduling primitive

Containers provide a runtime for a process

Use cases:

- Create
- Update
- Rollback
- Scale
- Pause/Resume

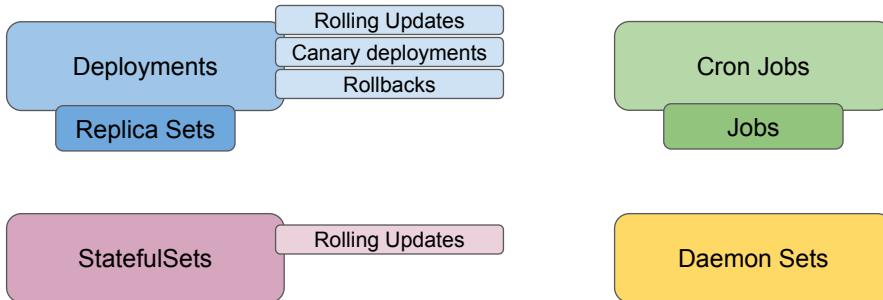
Exercise 1 - Second half



01-pods-deployments.md

- Create deployments/objects in your namespace
- Setup a simple nginx:1.7.9 deployment with single pod
- Try deleting the pod. Does it come up automatically?
- Can you access this pod from outside the cluster?

Controllers in Kubernetes



eficode

Kubernetes Fundamentals www.eficode.com

Service Discovery and Load Balancing



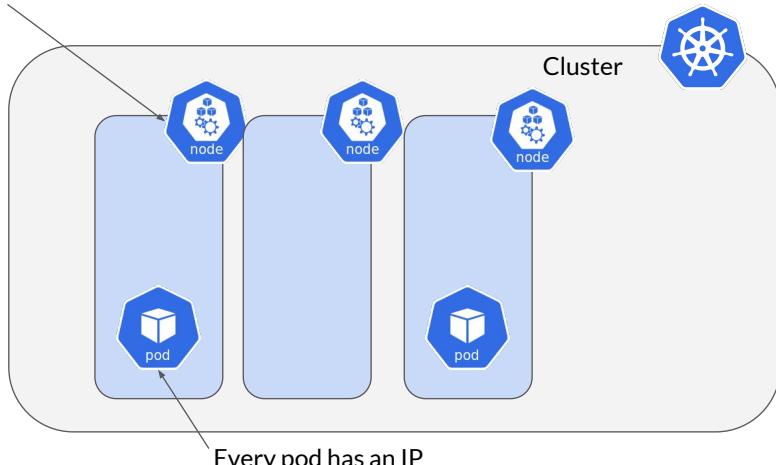
eficode

Kubernetes Fundamentals www.eficode.com

IP in Kubernetes



Every node has an IP



Service can have an IP

- Cluster
- NodePort
- Load Balancer



- Deployments and replica sets do not have an IP

Kubernetes Fundamentals www.eficode.com

eficode

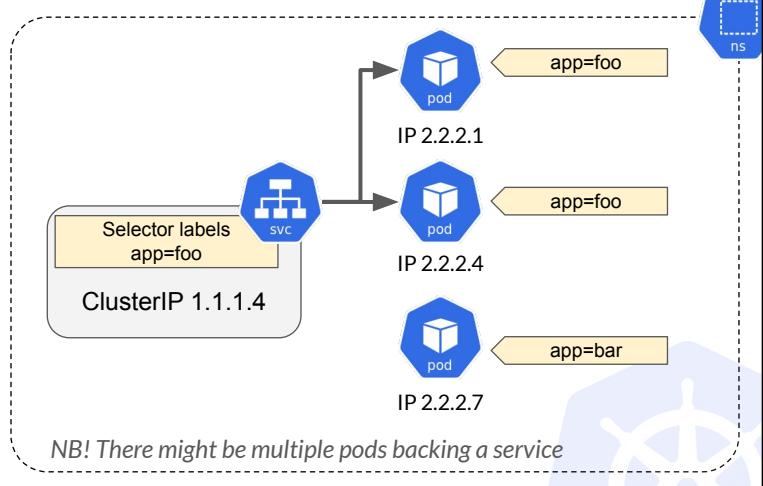
Service Discovery in Kubernetes



- Pods are ephemeral
 - IP's are dynamic
- Exposed through named Service
 - Selected through labels

Service can have an IP

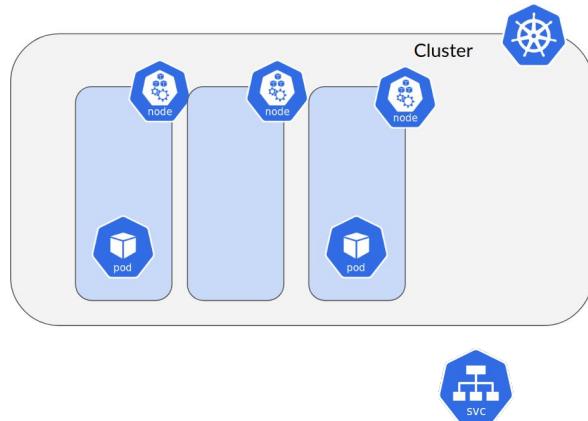
- Cluster
- NodePort
- Load Balancer



Kubernetes Fundamentals www.eficode.com

eficode

Service Discovery in Kubernetes



DNS names

Pod:

`<ip>. <namespace>. pod.cluster.local`

Service:

`<name>. <namespace>. svc.cluster.local`

eficode

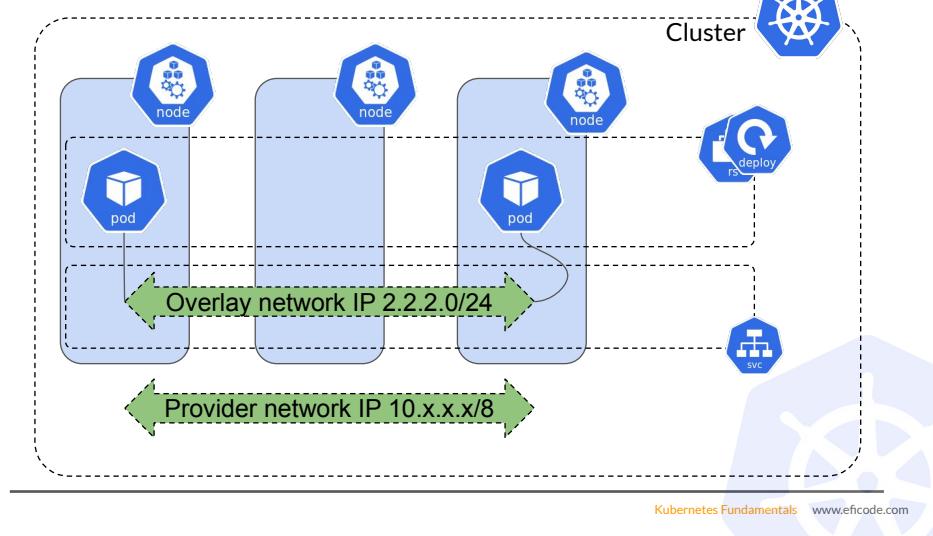
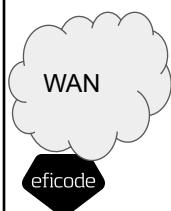
Kubernetes Fundamentals www.eficode.com

Service Discovery in Kubernetes



Cluster IP

- For internal communication
- The IP is static to the service



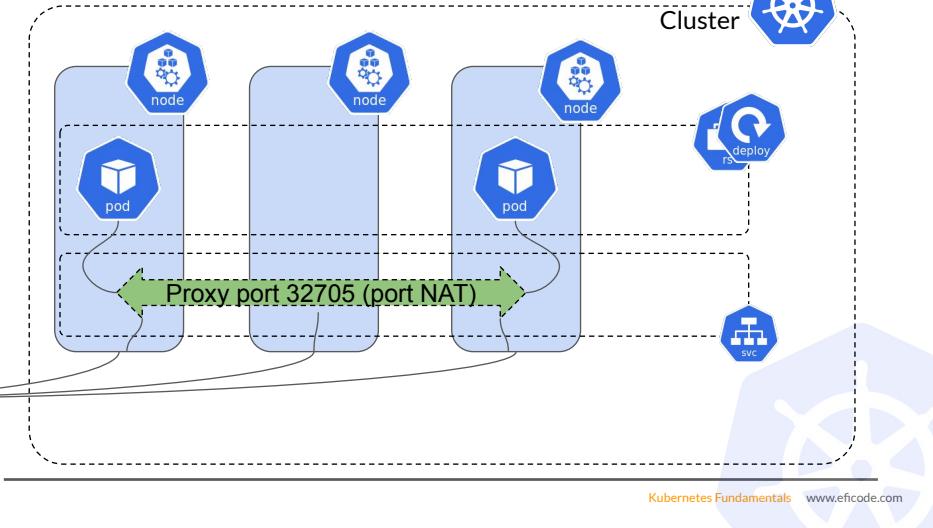
Kubernetes Fundamentals www.eficode.com

Service Discovery in Kubernetes



Nodeport

All nodes assigns the same port to a specific service



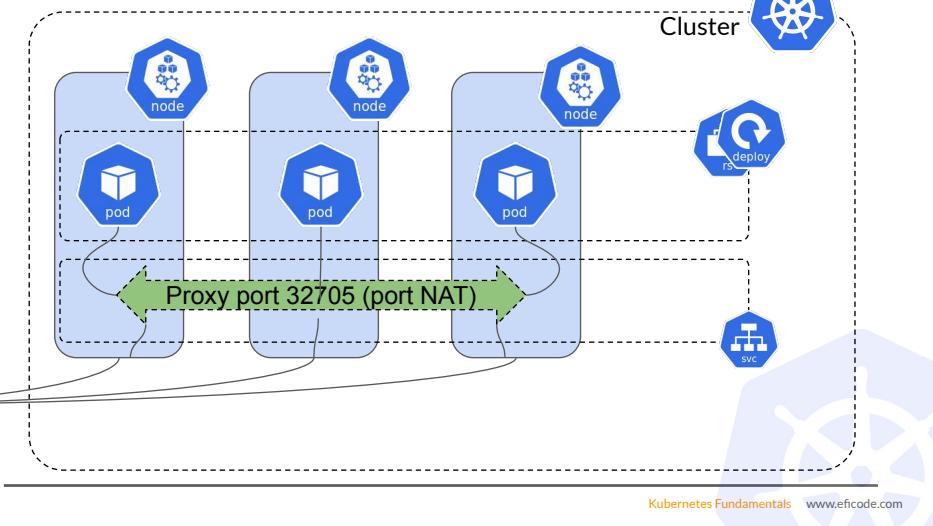
Kubernetes Fundamentals www.eficode.com

Service Discovery in Kubernetes



Nodeport

If a new pod is created, traffic is automatically routed to that pod as well

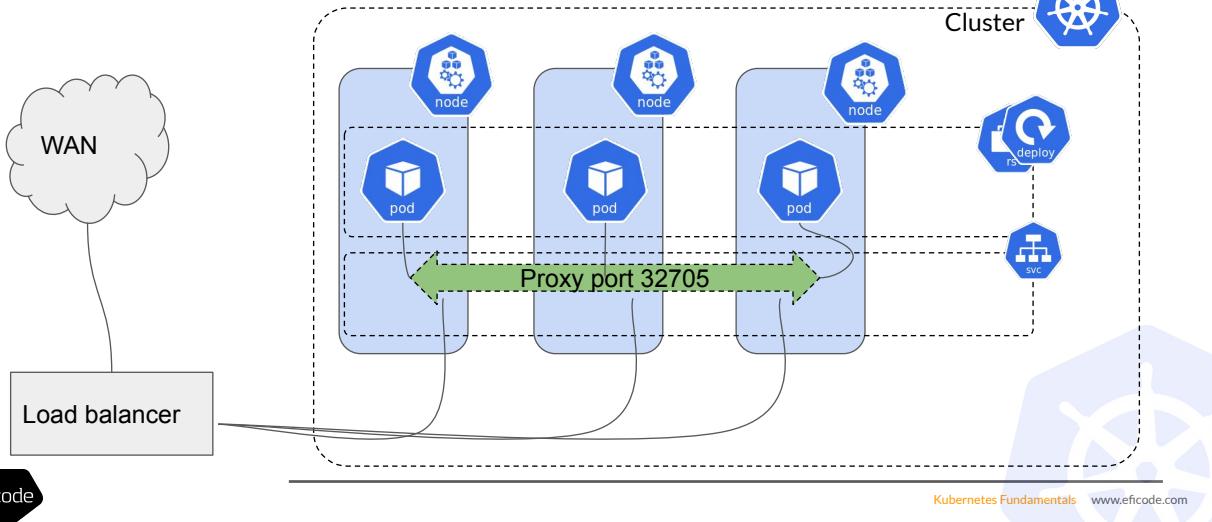


Kubernetes Fundamentals www.eficode.com

Service Discovery in Kubernetes



Load balancer



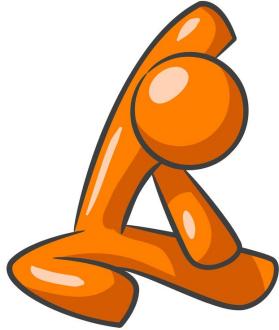
Network Abstraction - L4/Services

- Services are a networking L4 ‘virtual IP’ abstraction
 - Cluster IP address is fixed while backends can change as PODs change
 - Traffic is load-balanced (round-robin, least connections etc.)
- DNS names
 - Don’t use IP addresses in your applications, **always use DNS**
 - DNS search path reflects service name and namespace:

`<service-name>.<namespace>.svc.cluster.local`

In the same namespace we can use DNS short-name “<service-name>”

Exercise 2



02-service-discovery-loadbalancing.md

- Expose the deployment - as a service - type ClusterIP
- Can you access this service from outside the cluster?
- Delete / re-create the same service as type NodePort
- Can you access this service from outside the cluster?
- Delete / re-create the same service as type LoadBalancer
- Can you access this pod from outside the cluster?
- Increase number of replicas to four (4)
- What content you see each time you access the service?

Rolling out updates





Updating Resources

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 4
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

eficode

We can update/patch resources to change our application blueprint, e.g. replicas and image versions

Consider if this is best practice!

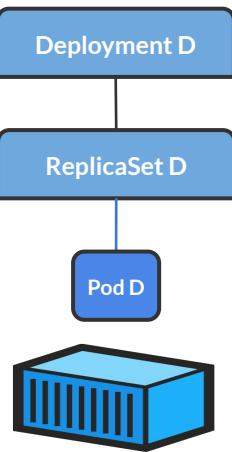
- \$ kubectl label ...
- \$ kubectl set <env | image | resources | ...> ...
- \$ kubectl edit ...
- \$ kubectl patch ...
- \$ kubectl scale --replicas 3 ...
- \$ kubectl apply -f manifest.yaml
- \$ kubectl replace ...

Kubernetes has versioned update for the **template** part of Deployments, DaemonSets and StatefulSets

- *All other changes are in-place!*

Kubernetes Fundamentals www.eficode.com

Managing Deployments in Kubernetes



Deployments manage the release of new versions of applications

ReplicaSets define how many instances should be running

Pods are the k8s scheduling primitive

Containers provide a runtime for a process



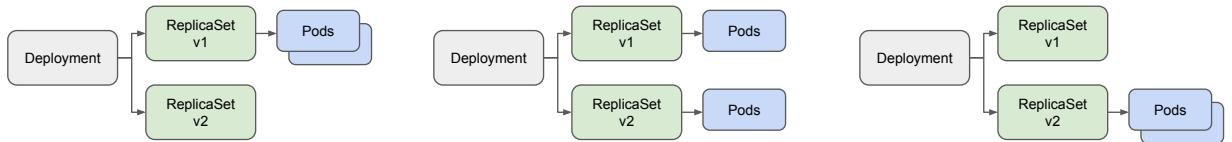
eficode

Kubernetes Fundamentals www.eficode.com

Rollout Strategies



All controllers have **RollingUpdate**



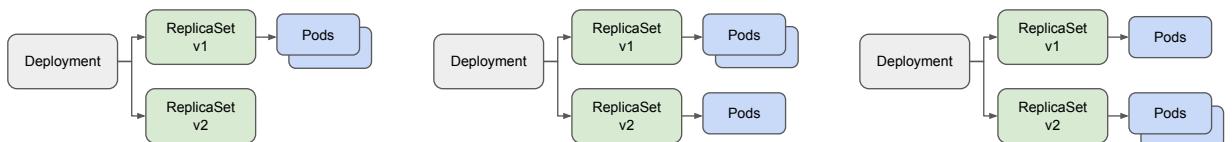
Rollout Strategies



RollingUpdate on Deployments have parameters:

- **maxSurge** (- number of pods above desired number)
- **maxUnavailable**
 - **maxSurge** and **maxUnavailable** can't be 0 at the same time

Example: maxSurge = 1, maxUnavailable = 0



Rollout History

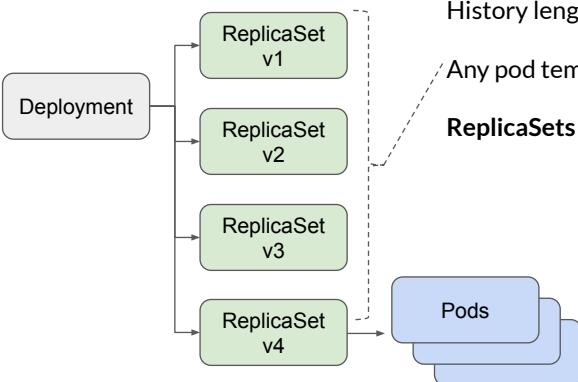


Deployments' history is tracked in ReplicaSets

History length is defined by `spec.revisionHistoryLimit`

/ Any pod template change triggers a new ReplicaSet / "new version"

ReplicaSets are named using the hash of the pod template



Rollout History Caveats (1)

```
$ kubectl create deployment nginx --image nginx:1.13.6
$ kubectl patch deployment nginx -p '{"spec": {"revisionHistoryLimit": 10}}'
$ kubectl set image deploy nginx nginx=nginx:1.13.7 --record
$ kubectl set image deploy nginx nginx=nginx:1.13.8 --record
$ kubectl set image deploy nginx nginx=nginx:1.13.9 --record
$ kubectl rollout history deploy nginx
REVISION  CHANGE-CAUSE
1          kubectl create deployment nginx --image=nginx:1.13.6 --record=true
2          kubectl set image deploy nginx nginx=nginx:1.13.7 --record=true
3          kubectl set image deploy nginx nginx=nginx:1.13.8 --record=true
4          kubectl set image deploy nginx nginx=nginx:1.13.9 --record=true
$ kubectl set image deploy nginx nginx=nginx:1.13.8 --record
$ kubectl rollout history deploy nginx
REVISION  CHANGE-CAUSE
1          kubectl create deployment nginx --image=nginx:1.13.6 --record=true
2          kubectl set image deploy nginx nginx=nginx:1.13.7 --record=true
4          kubectl set image deploy nginx nginx=nginx:1.13.9 --record=true
5          kubectl set image deploy nginx nginx=nginx:1.13.8 --record=true
```

This is what brings us trouble since we will be reusing an existing ReplicaSet

Ehh...? 1.13.7 is the third step in history, and revision 3 is missing

Expected, 1.13.8 is latest version and 1.13.9 is previous version

Rollout History Caveats (2)



```
$ kubectl rollout undo deployment nginx
$ kubectl rollout history deploy nginx
REVISION CHANGE-CAUSE
1   kubectl create deployment nginx --image=nginx:1.13.6 --record=true
2   kubectl set image deploy nginx nginx=nginx:1.13.7 --record=true
5   kubectl set image deploy nginx nginx=nginx:1.13.8 --record=true
6   kubectl set image deploy nginx nginx=nginx:1.13.9 --record=true

$ kubectl rollout undo deployment nginx
$ kubectl rollout history deploy nginx
REVISION CHANGE-CAUSE
1   kubectl create deployment nginx --image=nginx:1.13.6 --record=true
2   kubectl set image deploy nginx nginx=nginx:1.13.7 --record=true
6   kubectl set image deploy nginx nginx=nginx:1.13.9 --record=true
7   kubectl set image deploy nginx nginx=nginx:1.13.8 --record=true

$ kubectl rollout undo deployment nginx
$ kubectl rollout history deploy nginx
REVISION CHANGE-CAUSE
1   kubectl create deployment nginx --image=nginx:1.13.6 --record=true
2   kubectl set image deploy nginx nginx=nginx:1.13.7 --record=true
7   kubectl set image deploy nginx nginx=nginx:1.13.8 --record=true
8   kubectl set image deploy nginx nginx=nginx:1.13.9 --record=true
```

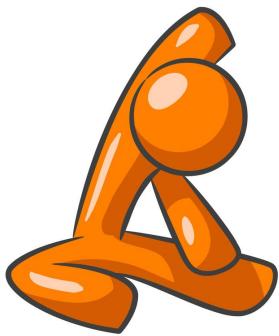
Another rollout undo
brings us back to 1.13.9

We never come back to
1.13.7

Exercise 3



03-rolling-updates.md



- Learn how to apply changes to running deployments
- Use the previous nginx:1.7.9 deployment with four replicas
- Rollout an update to nginx:1.9.1
- Test with **curl** from local computer
- Describe the deployment and replicaSets and notice the difference

Rollout Strategies - RollingUpdate

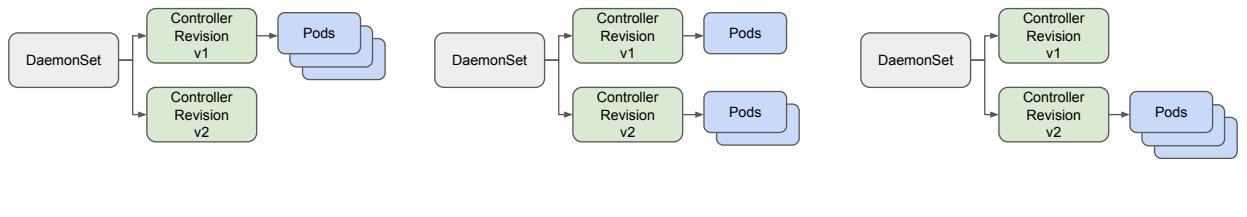


RollingUpdate on DaemonSets have parameters:

- **maxUnavailable**

- Can't be 0

Example, maxUnavailable = 2



Rollout Strategies - RollingUpdate

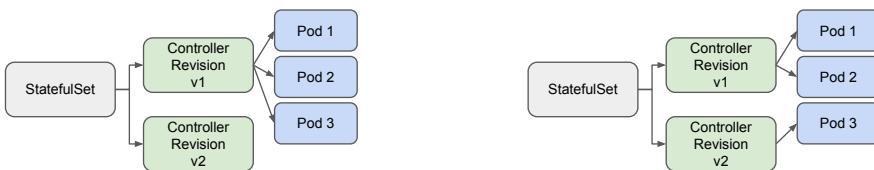


RollingUpdate on StatefulSets have parameters:

- **Partition**

- Only updates pods with "number" higher than **partition**
 - If **partition** is set higher than **replicas**, no pods are updated

Example, partition = 2



Rollout Strategies - RollingUpdate, more



spec.minReadySeconds

- Helpful for slowing down rolling updates
- Is set on the Deployment / StatefulSet / DaemonSet spec
 - Not on the **RollingUpdateStrategy** (just like the **revisionHistoryLimit**)

Other Rollout Strategies



Deployments have Recreate

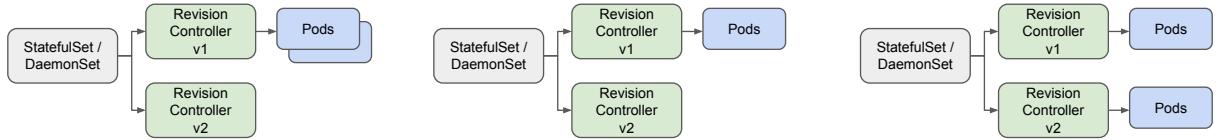


Other Rollout Strategies



StatefulSets and DaemonSets have OnDelete

(Manual deletion)



Break - Recap



- Pods, Deployments
- Redundancy
- Exposing with services (NodePort and LoadBalancer)
- Labels are important
- Updating replicas and images with different strategies
- Rollback

ConfigMaps and Secrets



12 Factor Applications



Methodology for application design - resonates well with the thinking behind Kubernetes

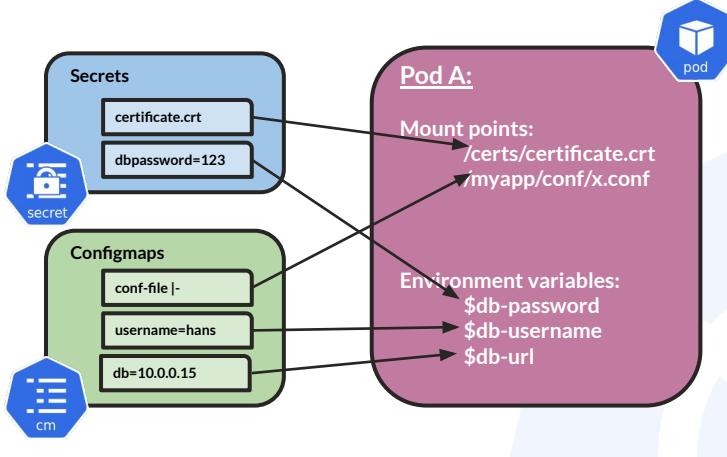
- Log to stderr/stdout
- Separate build, release and run stages
- Disposability and statelessness
- Horizontal scale-out
- Backing services are attached services
- Separate configuration and secrets from the executable

- Handle missing dependencies
 - Config files should be optional and have sane defaults
 - Retry access to external services - eliminate the need to deploy services in a specific order

Configuration Objects in Kubernetes



Configmaps and Secrets configure PODs



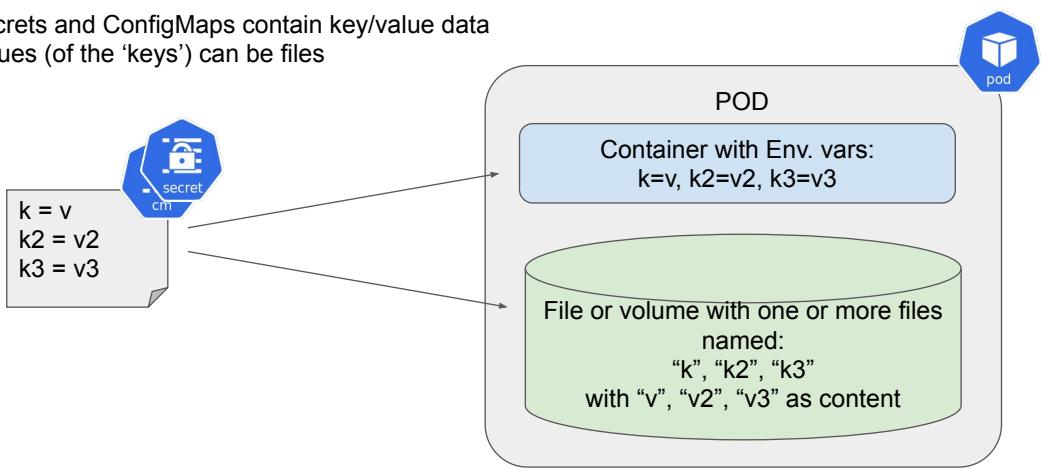
eficode

Kubernetes Fundamentals www.eficode.com

Configuration Objects in Kubernetes



- Secrets and ConfigMaps contain key/value data
- Values (of the 'keys') can be files



eficode

Kubernetes Fundamentals www.eficode.com

Configuration Objects in Kubernetes



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: language
data:
  LANGUAGE: Orcish
```

We inject configuration as environment variables

```
apiVersion: v1
kind: Secret
metadata:
  name: apikey
data:
  API_KEY: b25lcmluZ3R...
```

Using references to ConfigMaps and Secrets

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: envtest
...
spec:
  containers:
    - name: envtest
      ...
      env:
        - name: LANGUAGE
          valueFrom:
            configMapKeyRef:
              name: language
              key: LANGUAGE
        - name: API_KEY
          valueFrom:
            secretKeyRef:
              name: apikey
              key: API_KEY
```

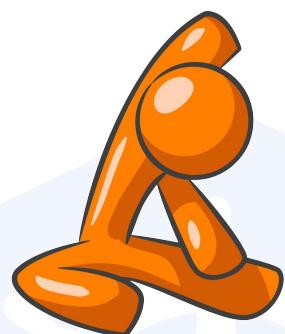
Kubernetes Fundamentals www.eficode.com

eficode

Exercise 4

04-secrets-configmap.md

- Overriding docker image environment variables with deployments



Kubernetes Fundamentals www.eficode.com

eficode

Automatic Pod Recreation (with Helm)



Automatically Roll Deployments When ConfigMaps or Secrets change

Often times configmaps or secrets are injected as configuration files in containers. Depending on the application a restart may be required should those be updated with a subsequent `helm upgrade`, but if the deployment spec itself didn't change the application keeps running with the old configuration resulting in an inconsistent deployment.

The `sha256sum` function can be used to ensure a deployment's annotation section is updated if another file changes:

```
kind: Deployment
spec:
  template:
    metadata:
      annotations:
        checksum/config: {{ include (print $.Template.BasePath "/configmap.yaml") . | sha256sum }}
```

Persistent Storage



Definition of Storage - Volumes



Why?

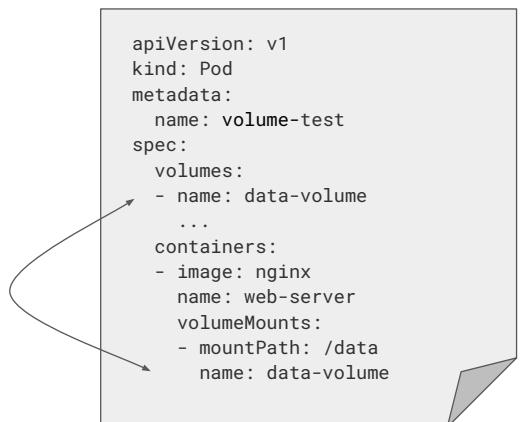
- Provide configuration for a container
- Persist data through pod lifecycle
- Share data between containers in a Pod

Mounting Volumes



How

- Volumes are specified in Pod-spec
- Mounted to one or more containers in their Container-spec



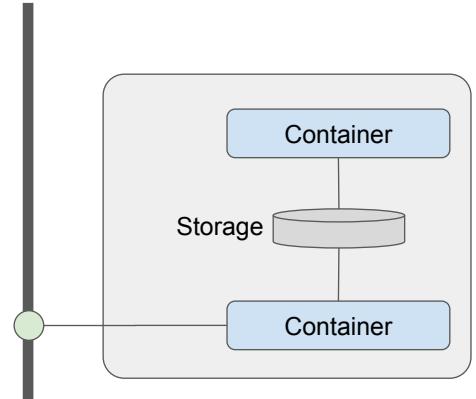
Sharing Data Between Containers in a Pod



Volume type `emptyDir`

- Only available during Pod lifecycle
- Great for sharing ephemeral data between containers in the same pod

```
...  
spec:  
  volumes:  
    - name: data-volume  
      emptyDir: {}
```



Mounting Volumes



Configmap

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: volume-test  
spec:  
  volumes:  
    - name: config-volume  
      configMap:  
        name: special-config  
        items:  
          - key: special.level  
            path: keys  
  containers:  
    - image: nginx  
      name: web-server  
      volumeMounts:  
        - mountPath: /config  
          name: config-volume
```

Volume Types



Generic and Platform Specific Types

- emptyDir, hostPath, local, nfs - generic types
- Secrets, ConfigMap, downwardAPI - accessing Kubernetes resources as volumes
- Cinder, cephfs etc. (many more) - network storage specific
- GCE, AWS, Azure disks - cloud provider integration
- PersistentVolumeClaim - dynamic provisioning

eficode

Kubernetes Fundamentals www.eficode.com

Dynamic Volume Provisioning



PersistentVolume and PersistentVolumeClaim - manual



Admin

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: hostpath
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  reclaimPolicy:
    - Recycle
  mountOptions:
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```



Dev / User

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: data-volume
spec:
  resources:
    requests:
      storage: 8Gi
  accessModes:
    - ReadWriteOnce
  storageClassName: ""
```

```
apiVersion: v1
kind: Pod
...
spec:
  volumes:
    - name: data-volume
      persistentVolumeClaim:
        claimName:
          data-volume
```

Kubernetes Fundamentals www.eficode.com

Dynamic Volume Provisioning

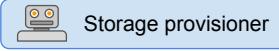


StorageClass and PersistentVolumeClaim - automatic



Admin

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: standard
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
```



Automatically provisioned

```
apiVersion: v1
kind: PersistentVolume
...
spec:
  gcePersistentDisk:
    fsType: ext4
    pdName: gke-cluster-1-3e90xx
```



Dev / User

```
apiVersion: v1
kind: Pod
...
spec:
  volumes:
    - name: data-volume
      persistentVolumeClaim:
        claimName: standard
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: data-volume
spec:
  resources:
    requests:
      storage: 8Gi
  accessModes:
    - ReadWriteOnce
  storageClassName: "standard"
```

eficode

Storage objects in Kubernetes



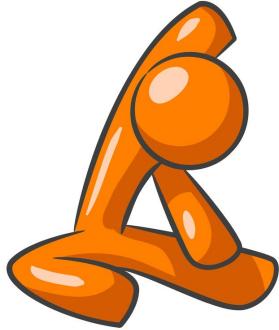
Storage in the cluster with objects:

Pods can reference a Volume of type PersistentVolumeClaim to acquire access to the storage

- Can be mounted to one or more nodes, with different modes
 - ReadWriteOnce - Mounted as **readwrite** on one **Node**
 - ReadOnlyMany - Mounted as **read** on many **Nodes**
 - ReadWriteMany - Mounted as **readwrite** on many **Nodes**
 - RWM is generally not supported by Volume Plugins

eficode

Exercise 5



05-storage.md

- Use the previous nginx:1.9.1 deployment
- Create a PVC using the instructions.
- Update nginx deployment to use this PVC as a mount point for its html directory.
- Deploy this change and test it with curl.

Helpful Task: <https://kubernetes.io> -> Documentation -> Tasks -> Configure Pods and Containers -> [Configure a Pod to Use a PersistentVolume for Storage](#)

Networking and Ingress



Network Abstraction - L4/Services

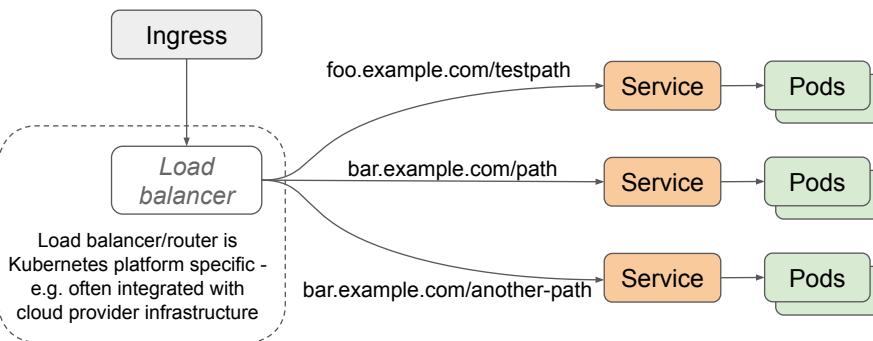


- Services are a networking L4 ‘virtual IP’ abstraction
 - Cluster IP address is fixed while backends can change as number of PODs change
 - Traffic is load balanced with several algorithms possible (round-robin, least connections etc.)
- Services are exposed through DNS
 - Don’t use IP addresses in your applications, **always use DNS**
 - DNS search path reflects service name and namespace
 - <service-name>.<namespace>.svc.cluster.local
 - Access your services using short DNS name <service-name>

Network Abstraction - L7/Ingress



- Ingress provides a networking L7 HTTP abstraction
- Defines load balancer/routing of L7 traffic

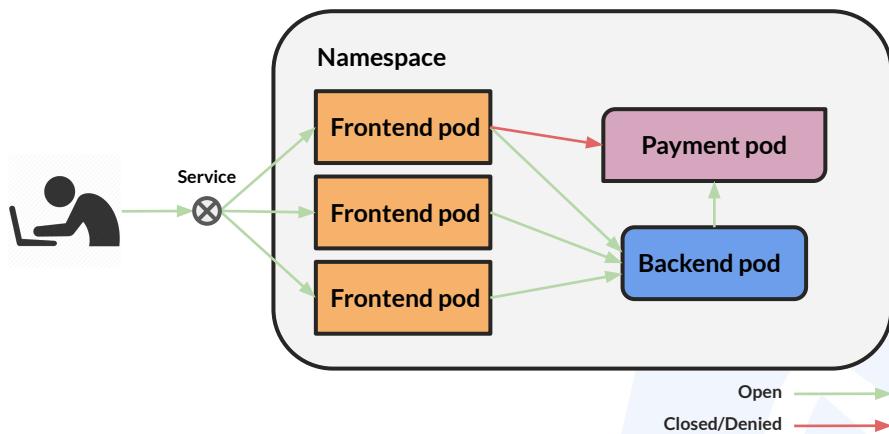


```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: web
spec:
  rules:
  - host: foo.example.com
    http:
      paths:
      - path: /testpath
        backend:
          serviceName: web
          servicePort: http
```

Security



Enforcing network policies



NetworkPolicy



```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: payment-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: payment-gateway
  policyTypes:
    - Ingress
    - Egress
...
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              project: webshop
        podSelector:
            matchLabels:
              app: backend
      ports:
        - protocol: TCP
          port: 1111
  egress:
    - to:
        - ipBlock:
            cidr: 1.2.3.4/32
        - ipBlock:
            cidr: 5.6.7.8/32
      ports:
        - protocol: TCP
          port: 9999
```

eficode

Partials www.eficode.com

POD Security Policies



Enable fine-grained authorization of POD features

- Use of host namespaces (PID, IPC, networking/host-ports)
- Use of volume types (including host path)
- Run container as root vs. non-root
- Read-only filesystem in pod
- Container capabilities (e.g. NET_RAW allows capture packets)
- Privileged container (e.g. low-level access to host devices through /dev)

eficode

Kubernetes Fundamentals www.eficode.com

POD Security Policies



```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
spec:
  allowedCapabilities:
    - SYS_ADMIN
    - NET_ADMIN
  allowedHostPaths:
    - pathPrefix: /tmp
  fsGroup:
    rule: RunAsAny
  hostIPC: false
  hostNetwork: false
  volumes:
    - hostPath
    - configMap
    - secret
```

- PSPs are Cluster-scoped - no namespaces
- PSP support needs to be enabled in the cluster
 - One of the many Kubernetes admission-controllers
- Allow-approach: **Default is DENY**
 - There must be a policy that ALLOW use of feature

POD Security Policies



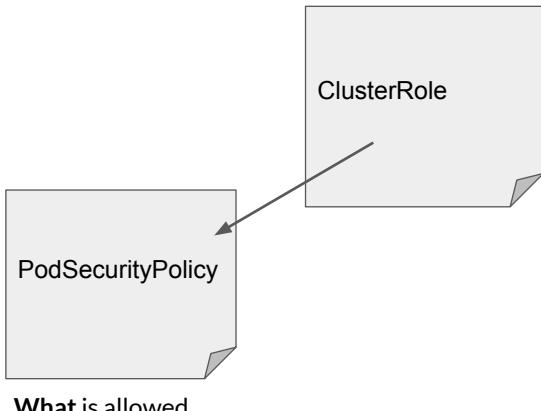
PodSecurityPolicy

What is allowed...

POD Security Policies



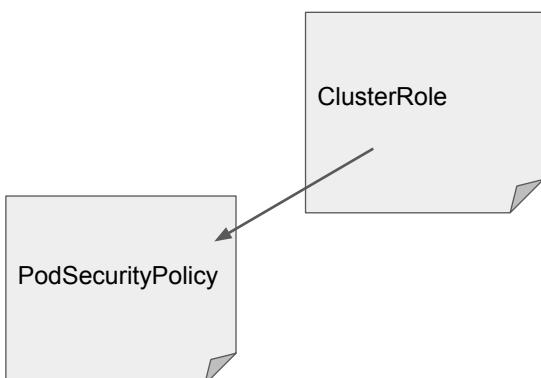
Role to allow **use** of policy...



POD Security Policies



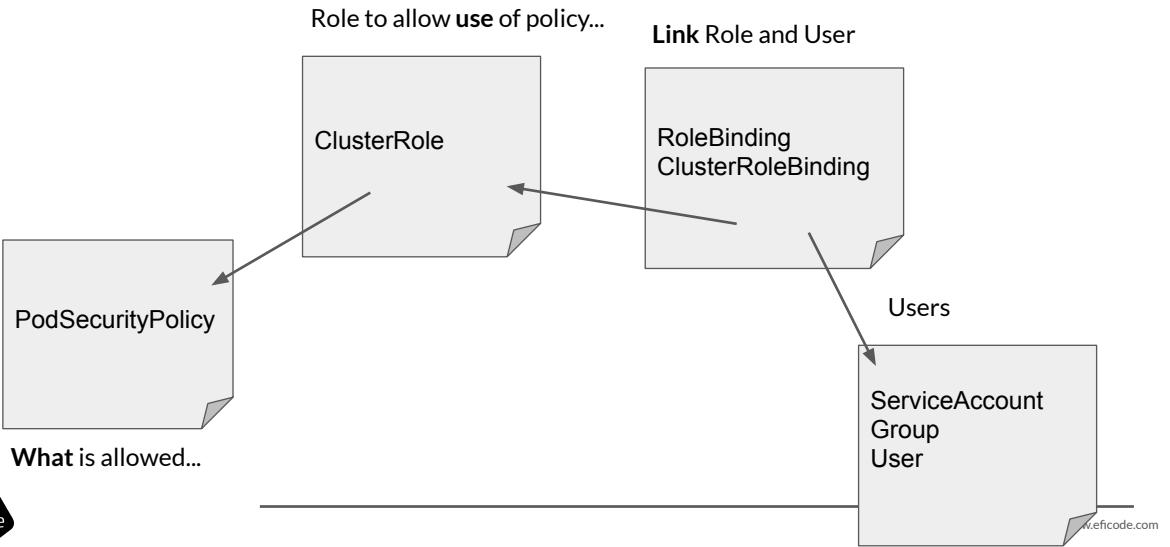
Role to allow **use** of policy...



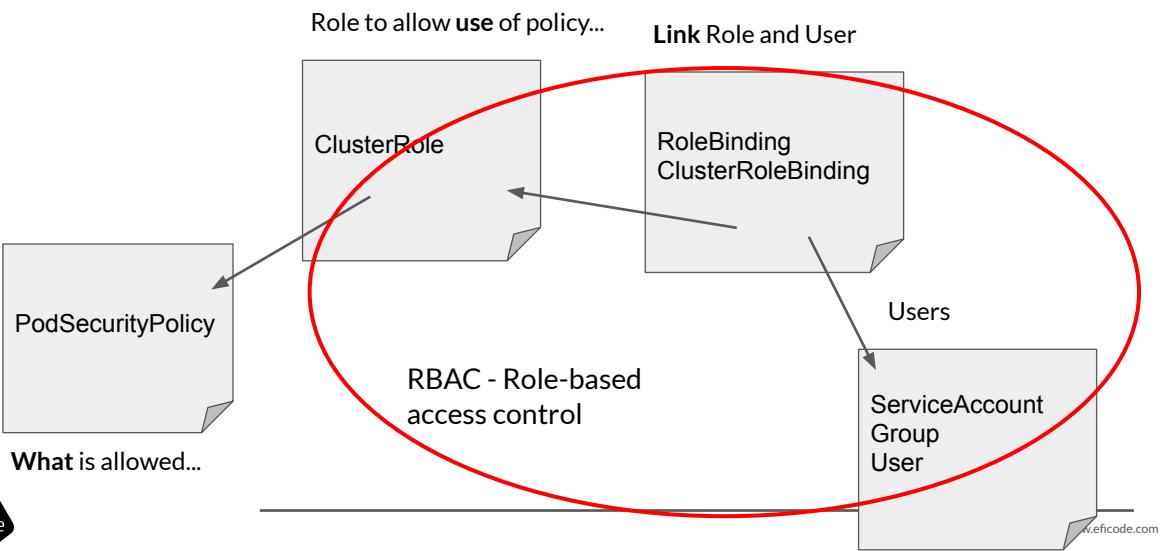
Users



POD Security Policies



POD Security Policies



Liveness and Health Checks



Readiness and Health Checks

- Defaults
 - No probe in pod spec imply ready/alive
 - Init-containers imply not ready
- Essential part of the ‘declarative approach’

```
...  
spec:  
  containers:  
    - name: my-app  
      image: foo  
      livenessProbe:  
        httpGet:  
          path: /  
          Port: 1234  
      initialDelaySeconds: 5  
      periodSeconds: 5
```

```
...  
spec:  
  containers:  
    - name: my-app  
      image: foo  
      readinessProbe:  
        exec:  
          command:  
            - cat  
            - /tmp/ready  
      initialDelaySeconds: 5  
      periodSeconds: 5
```

Readiness and Health Checks



- Liveness probing
 - Container restart depends on health checks
- Readiness probing
 - Pods must be ready before being added as service backends
 - Important in statefulsets scaling
- Probing types
 - HTTP (Return 2xx or 3xx, no authentication)
 - TCP (3-way-handshake indicates success)
 - Exec (return code 0)

```
...  
spec:  
  containers:  
    - name: my-app  
      image: foo  
      livenessProbe:  
        httpGet:  
          path: /  
          port: 1234  
      initialDelaySeconds: 5  
      periodSeconds: 5
```

```
...  
spec:  
  containers:  
    - name: my-app  
      image: foo  
      readinessProbe:  
        exec:  
          command:  
            - cat  
            - /tmp/ready  
      initialDelaySeconds: 5  
      periodSeconds: 5
```

eficode

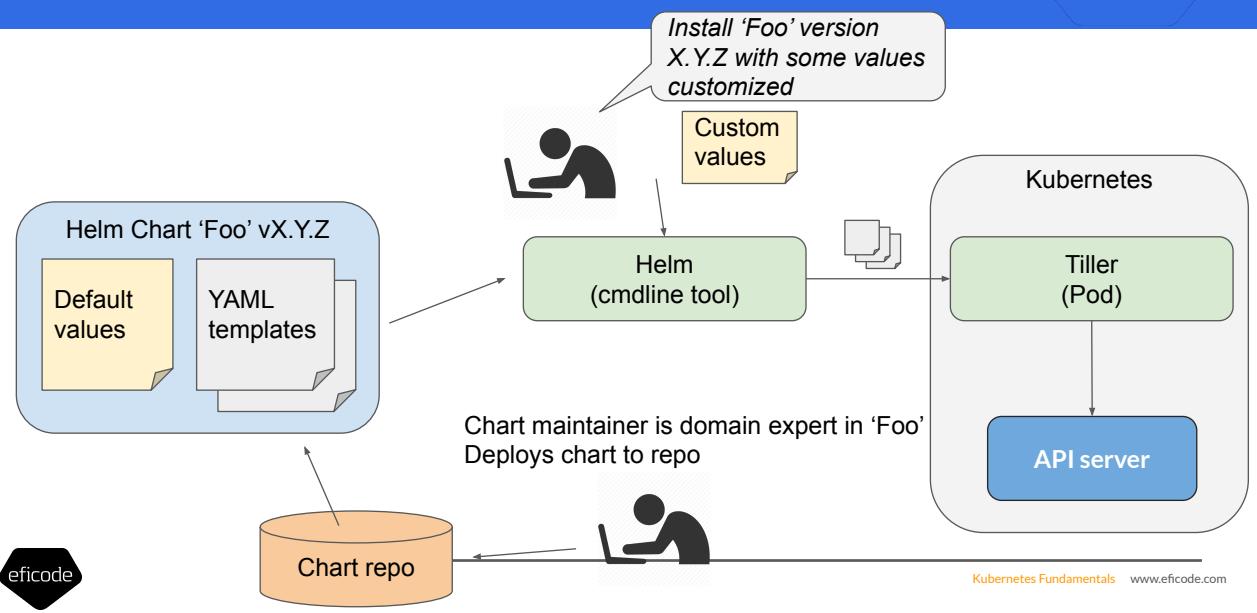
Kubernetes Package Management with Helm



eficode

Kubernetes Fundamentals www.eficode.com

Helm Package Manager

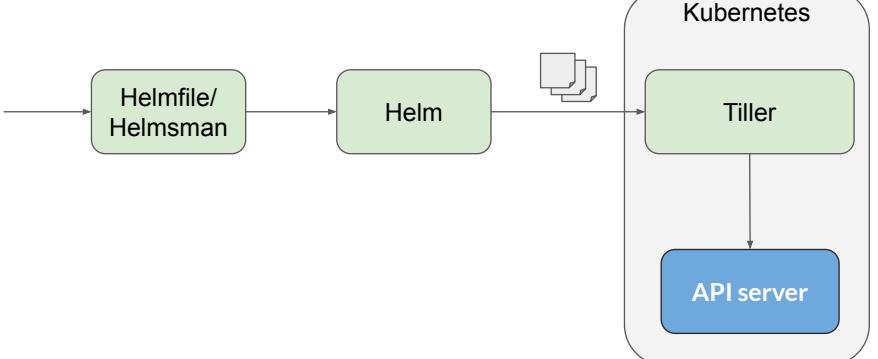


Managing Apps With Helmfile/Helmsman

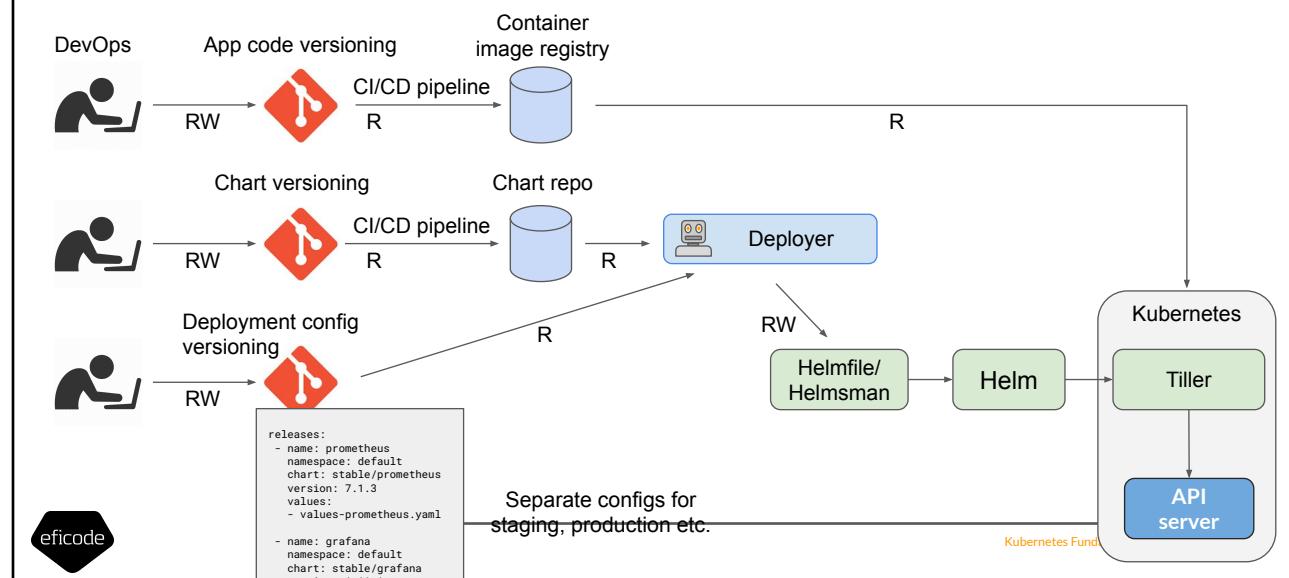


```
releases:
- name: prometheus
  namespace: default
  chart: stable/prometheus
  version: 7.1.3
  values:
    - values-prometheus.yaml

- name: grafana
  namespace: default
  chart: stable/grafana
  version: 1.11.1
  values:
    - values-grafana.yaml
```



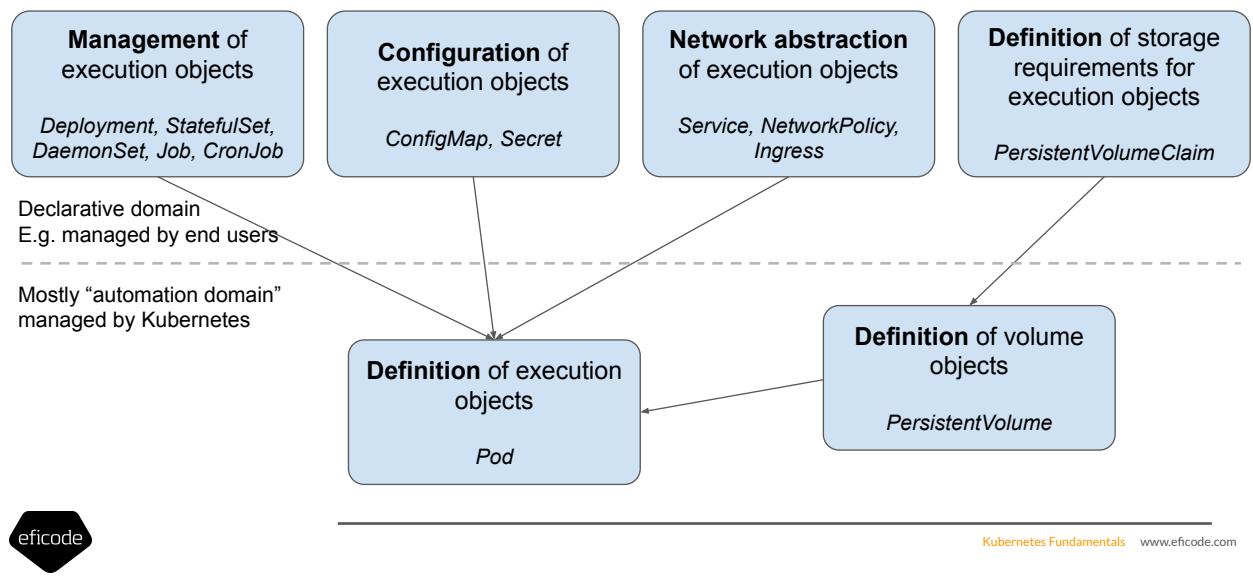
Managing Apps - the Big Picture



Wrap-up / Extras



Remember the Kubernetes Lego Bricks?



Playing with K8s



- www.k8s.io - Everything Kubernetes / learning
- www.KataCoda.com - online sandbox / learning
- www.k3s.io - minimal installation

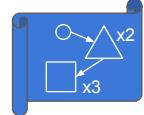
Kubernetes Internals



The Short Description of Kubernetes



Application blueprint

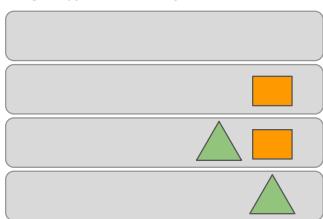


*"Build and maintain
my application
according to this
blueprint"*

Blueprint storage



Container nodes



Controllers

Scheduling, resource mgr

Container instance manager

Storage manager

Security compliance mgr

Network manager

... etc

The Short Description of Kubernetes (2)



Application blueprint



Application blueprint defines application structure using building blocks known as **Kubernetes resources**

Application architecture is declared to the **Kubernetes API server**

HA blueprint storage is an essential concept in Kubernetes (implemented by **etcd**)

Container nodes

Containers

Containers are scheduled to **worker nodes** and maintained, e.g. restarted in case nodes fail

Blueprint storage



Kubernetes automation is handled by controllers (mostly on the **Kubernetes master nodes**)

Controllers

Scheduling, resource mgr

Container instance manager

Storage manager

Security compliance mgr

Network manager

Kubernetes controllers continuously compare the state of the application with the blueprint and automatically correct differences

eficode

Components of Kubernetes



Master(s)

Controller manager

Scheduler



Admin / Dev / User

Kube-proxy

Kubelet

Kube-proxy

Kubelet

Pod A

Pod B

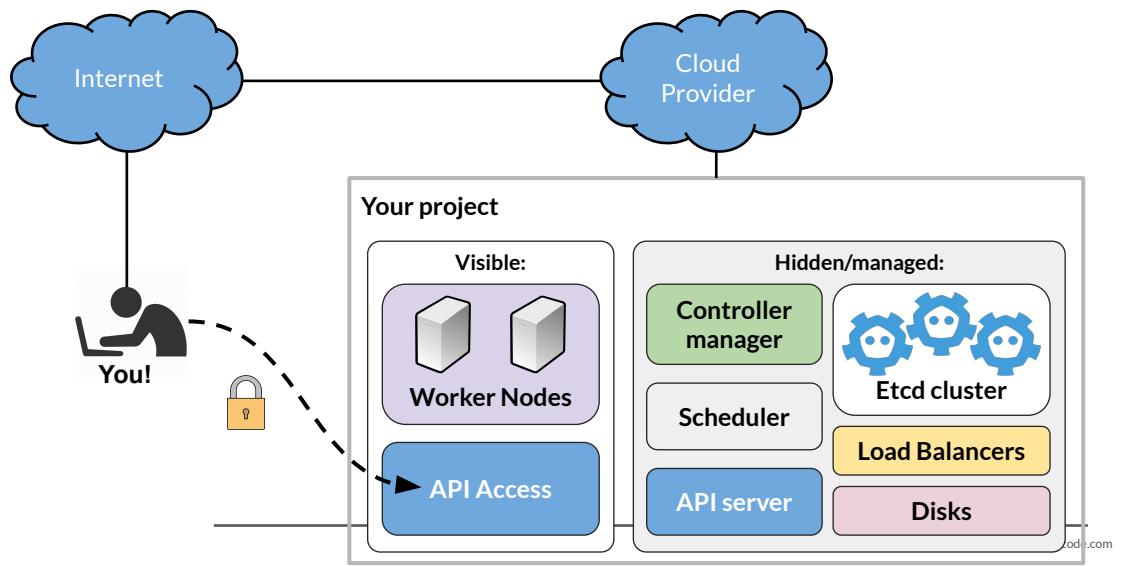
Pod D

Pod E

Pod C

eficode

Cluster layout: K8s-aaS (GCP, AWS, ...)



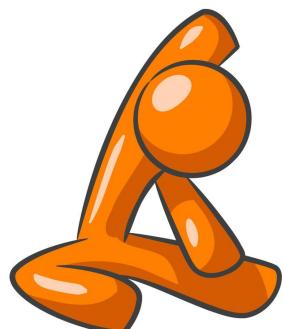
eficode

.com

Agenda



- Introduction to Kubernetes
- Namespaces, Pods, and Deployments
- Service Discovery and LoadBalancing
- Rolling out updates
- ConfigMaps and Secrets
- Persistent Storage
- Networking and Ingress
- Helm Package Manager
- Kubernetes Internals



eficode



eficode



eficode

Kubernetes Fundamentals www.eficode.com